



```
1 student@student-HP-Pro-3330-MT:~$ sudo su hduser
2 [sudo] password for student:
3 hduser@student-HP-Pro-3330-MT:/home/student$ cd
4 hduser@student-HP-Pro-3330-MT:~/ start-dfs.sh
5 25/04/07 13:28:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...
   using builtin-java classes where applicable
6 Starting namenodes on [localhost]
7 localhost: starting namenode, logging to /usr/local/hadoop/logs/hadoop-hduser-namenode-student-HP-
   Pro-3330-MT.out
8 localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-hduser-datanode-student-HP-
   Pro-3330-MT.out
9 localhost: OpenJDK Server VM warning: You have loaded library /usr/local/hadoop/lib/native/
   libhadoop.so.1.0.0 which might have disabled stack guard. The VM will try to fix the stack guard now.
10 localhost: It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it
    with '-z noexecstack'.
11 Starting secondary namenodes [0.0.0.0]
12 0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop/logs/hadoop-hduser-secondarynamenode-
   student-HP-Pro-3330-MT.out
13 0.0.0.0: OpenJDK Server VM warning: You have loaded library /usr/local/hadoop/lib/native/
   libhadoop.so.1.0.0 which might have disabled stack guard. The VM will try to fix the stack guard now.
14 0.0.0.0: It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with
    '-z noexecstack'.
15 25/04/07 13:28:50 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...
   using builtin-java classes where applicable
16 hduser@student-HP-Pro-3330-MT:~$ start-yarn.sh
17 starting yarn daemons
18 starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-hduser-resourcemanager-student-HP-
```



Auto Menu - + INS

φ

Plain Text Tab Width: 8 Ln 298, Col 1 HP v185e

```
3 (~/Desktop) - gedit 2:10 PM
File Open Save Undo Redo Cut Copy Paste Find Replace
f1.txt x employee.txt x emp.txt x prac3 x
19 localhost: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-hduser-nodemanager-student-HP-Pro-3330-MT.out
20 localhost: OpenJDK Server VM warning: You have loaded library /usr/local/hadoop/lib/native/libhadoop.so.1.0.0 which might have disabled stack guard. The VM will try to fix the stack guard now.
21 localhost: It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with '-z noexecstack'.
22 hduser@student-HP-Pro-3330-MT:~$ jps
23 10230 NameNode
24 10609 SecondaryNameNode
25 10400 DataNode
26 10913 NodeManager
27 10777 ResourceManager
28 11269 Jps
29 hduser@student-HP-Pro-3330-MT:~$ cd
30 hduser@student-HP-Pro-3330-MT:~$ cd /usr/local/hive/bin
31 hduser@student-HP-Pro-3330-MT:/usr/local/hive/bin$ ./hive
32 SLF4J: Class path contains multiple SLF4J bindings.
33 SLF4J: Found binding in [jar:file:/usr/local/hive/lib/log4j-slf4j-impl-2.4.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
34 SLF4J: Found binding in [jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
35 SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
36 SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
37
38 Logging initialized using configuration in jar:file:/usr/local/hive/lib/hive-common-2.1.1.jar!/hive-log4j2.properties Async: true
39 Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different configuration for your use case.
```

Plain Text Tab Width: 8 Ln 298, Col 1 INS



Auto Menu

HP v185e

prac3 (~/Desktop) - gedit

```
log4j2.properties Async: true
39 Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a
different execution engine (i.e. tez, spark) or using Hive 1.X releases.
40 hive> show databases;
41 OK
42 default
43 flightdb
44 hive
45 jyoti
46 kc
47 sanket
48 Time taken: 0.771 seconds, Fetched: 6 row(s)
49 hive> create database if not exists new_jyoti;
50 FAILED: ParseException line 1:23 missing KW_EXISTS at 'exists' near '<EOF>'
51 line 1:29 extraneous input 'new_jyoti' expecting EOF near '<EOF>' I
52 hive> create database if not exists new_jyoti;
53 OK
54 Time taken: 0.349 seconds
55 hive> use new_jyoti;
56 OK
57 Time taken: 0.015 seconds
58 hive> create table if not exists emps(id int, name string) row format delimited fields terminated by ','
      stored as textfile;
59 OK
60 Time taken: 0.634 seconds
61 hive> show tables;
62 OK
--
```

Plain Text — TabWidth: 8 — Ln 298, Col 1

INS

HP v185e

```
f1.txt x employee.txt x emp.txt x prac3 x
61 hive> show tables;
62 OK
63 emps
64 Time taken: 0.049 seconds, Fetched: 1 row(s)
65 hive> describe tables;
66 FAILED: SemanticException [Error 10001]: Table not found tables
67 hive> describe emps;
68 OK
69 id          int
70 name        string
71 Time taken: 0.375 seconds, Fetched: 2 row(s)
72 hive> load data local inpath '/home/student/Desktop/emp.txt' into table emps;
73 Loading data to table new_jyoti.emps
74 OK
75 Time taken: 1.085 seconds
76 hive> select *From emps;
77 OK
78 1      jyoti
79 2      tau
80 3      neha
81 4      kalyani
82 NULL    NULL
83 Time taken: 1.195 seconds, Fetched: 5 row(s)
84 hive>
85 hive> create table if not exists emps(id int, name string) row format delimited fields terminated by ','
         stored as textfile;
86 OK
87 Time taken: 0.634 seconds
```

Plain Text \* Tab Width: 8 \* Ln 298, Col 1 INS

HP v185e

prac3 (~/Desktop) - gedit

```
88hive> show tables;
89 OK
90 emps
91 Time taken: 0.049 seconds, Fetched: 1 row(s)
92hive> describe tables;
93 FAILED: SemanticException [Error 10001]: Table not found tables
94hive> describe emps;
95 OK
96 id          int
97 name        string
98 Time taken: 0.375 seconds, Fetched: 2 row(s)
99hive> load data local inpath '/home/student/Desktop/emp.txt' into table emps;
100 Loading data to table new_jyoti.emps
101 OK
102 Time taken: 1.085 seconds
103hive> select *From emps;
104 OK
105 1      jyoti
106 2      tau
107 3      neha
108 4      kalyani
109 NULL    NULL
110 Time taken: 1.195 seconds, Fetched: 5 row(s)
111hive>
112hive> set hive.txn.manager=org.apache.hadoop.hive ql.lockmgr.DbTxnManager;
113hive> set hive.support.concurrency=true;
114hive> update emps set id = 5 where name='Jyoti';
115 FAILED: SemanticException [Error 102971]: Attempt to do update or delete on table new_jyoti.emps that
```

Plain Text Tab Width: 8 Ln 298, Col 1 INS



Auto Menu

HP vi

```
117 hive> set hive.support.concurrency=true;
118 hive> update emps set id = 5 where name='Jyoti';
119 FAILED: SemanticException [Error 10297]: Attempt to do update or delete on table new_jyoti.emps that
120 does not use an AcidOutputFormat or is not bucketed
121 hive> alter table emps rename to emp_new;
122 OK
123 Time taken: 0.447 seconds
124 hive> alter table emp add columns (age int);
125 FAILED: SemanticException [Error 10001]: Table not found new_jyoti.emp
126 hive> alter table emp_new add columns (age int);
127 OK
128 Time taken: 0.212 seconds
129 hive> create index idx_emp_name on table emp_new(name) as
130   'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler' with deferred rebuild;
131 OK
132 Time taken: 0.434 seconds
133 hive> alter index idx_emp_name on emp_new rebuild;
134 WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider
135   using a different execution engine (i.e. tez, spark) or using Hive 1.X releases.
136 Query ID = hduser_20250407134623_cc51f4fd-195f-4947-a927-522f9b4d5b4a
137 Total jobs = 1
138 Launching Job 1 out of 1
139 Number of reduce tasks not specified. Estimated from input data size: 1
140 In order to change the average load for a reducer (in bytes):
141   set hive.exec.reducers.bytes.per.reducer=<number>
142 In order to limit the maximum number of reducers:
```

Plain Text Tab Width: 8 Ln 298, Col 1 INS

HP v185e

```
f1.txt x employee.txt x emp.txt x prac3 x
In order to change the average load for a reducer (in bytes).
535 set hive.exec.reducers.bytes.per.reducer=<number>
36 In order to limit the maximum number of reducers:
37 set hive.exec.reducers.max=<number>
38 In order to set a constant number of reducers:
39 set mapreduce.job.reduces=<number>
40 Starting Job = job_1744012749162_0001, Tracking URL = http://student-HP-Pro-3330-MT:8088/proxy/
41 Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1744012749162_0001
42 Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
43 2025-04-07 13:46:47,375 Stage-1 map = 0%, reduce = 0%
44 Interrupting... Be patient, this might take some time.
45 Press Ctrl+C again to kill JVM
46 killing job with: job_1744012749162_0001
47 2025-04-07 13:47:02,312 Stage-1 map = 0%, reduce = 100%
48 2025-04-07 13:47:03,335 Stage-1 map = 100%, reduce = 100%
49 Ended Job = job_1744012749162_0001 with errors
50 Error during job, obtaining debugging information...
51 Examining task ID: task_1744012749162_0001_m_000000 (and more) from job job_1744012749162_0001
52
53 Task with the most failures(1):
54 -----
55 Task ID:
56 task_1744012749162_0001_m_000000
57
58 URL:
59 http://0.0.0.0:8088/taskdetails.jsp?jobid=job_1744012749162_0001&tipid=task_1744012749162_0001_m_000000
60 -----
```

Plain Text Tab Width: 8 Ln 298, Col 1 INS



Auto Menu

HP.v185

ektop) - gedit KKWECEIT/43/133

Open Save Undo Redo Cut Copy Paste Find Replace Find Next Find Previous Find in Files Find in Directories Find in Selection Find in All Files Find in All Directories Find in All Selections Find in All Files and Directories Find in All Selections and Directories Find in All Files and Directories and Selections Find in All Selections and Directories and Selections Find in All Files and Directories and Selections and Selections Find in All Selections and Directories and Selections and Selections Find in All Files and Directories and Selections and Selections and Selections Find in All Selections and Directories and Selections and Selections and Selections Find in All Files and Directories and Selections and Selections and Selections and Selections

f1.txt x employee.txt x emp.txt x prac3 x

```
59 http://0.0.0.0:8088/taskdetails.jsp?jobid=job_1744012749162_0001&tipid=task_1744012749162_0001_m_000000
60 -----
61 Diagnostic Messages for this Task:
62 Task KILL is received. Killing attempt!
63 [2025-04-07 13:47:02.404]Container killed by the ApplicationMaster.[2025-04-07 13:47:02.404]
64 [2025-04-07 13:47:03.441]Container killed on request. Exit code is 143[2025-04-07 13:47:03.441]
65 [2025-04-07 13:47:03.442]Container exited with a non-zero exit code 143. [2025-04-07 13:47:03.442]
66
67
168 FAILED: Execution Error, return code 2 from org.apache.hadoop.hive.ql.exec.mr.MapRedTask
169 MapReduce Jobs Launched:
170 Stage-Stage-1: Map: 1 Reduce: 1 HDFS Read: 0 HDFS Write: 0 FAIL
171 Total MapReduce CPU Time Spent: 0 msec
172 hive> show indexes on emp_new;
173 OK
174 idx_emp_name          emp_new >           name
   new_jyoti_emp_new_idx_emp_name_ compact
175 Time taken: 0.135 seconds, Fetched: 1 row(s)
176 hive> select *from Display all 528 possibilities? (y or n)
177 hive> select *from ew_jyoti_emp_new_idx_emp_name_
178   > select *from ew_jyoti_emp_new_idx_emp_name_;
179 FAILED: ParseException line 2:0 missing EOF at 'select' near 'ew_jyoti_emp_new_idx_emp_name_'
180 hive> select *from ew_jyoti_emp_new_idx_emp_name_;
181 FAILED: SemanticException [Error 10001]: Line 1:13 Table not found 'ew_jyoti_emp_new_idx_emp_name_'
182 hive> select *from new_jyoti_emp_new_idx_emp_name_;
183 OK
184 Time taken: 0.245 seconds
```

Plain Text Tab Width: 8 Ln 298, Col 1 INS

HP v185e

(~/Desktop) - gedit

```
183 OK
184 Time taken: 0.245 seconds
185 hive> show tables;
186 OK
187 emp_new
188 new_jyoti_emp_new_idx_emp_name_
189 Time taken: 0.093 seconds, Fetched: 2 row(s)
190 hive> select *from new_jyoti_emp_new_idx_emp_name_;
191 OK
192 Time taken: 0.227 seconds
193 hive>
194 hive> create index idx_emp_id on table emp_new(id) as
  'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler' with deferred rebuild;
195 OK
196 Time taken: 0.317 seconds
197 hive> alter index idx_emp_id on emp_new rebuild;
198 WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider
  using a different execution engine (i.e. tez, spark) or using Hive 1.X releases.
199 Query ID = hduser_20250407135230_0084ceda-3546-4eef-aed2-d408f4d08916
200 Total jobs = 1
201 Launching Job 1 out of 1
202 Number of reduce tasks not specified. Estimated from input data size: 1
203 In order to change the average load for a reducer (in bytes):
204   set hive.exec.reducers.bytes.per.reducer=<number>
205 In order to limit the maximum number of reducers:
206   set hive.exec.reducers.max=<number>
207 In order to set a constant number of reducers:
```

Plain Text Tab Width: 8 Ln 298, Col 1  
INS  
HP v185e

c3 (~/Desktop) - gedit

KKWEC/IT/43

```
206 set hive.exec.reducers.max=<number>
207 In order to set a constant number of reducers:
208 set mapreduce.job.reduces=<number>
209 Starting Job = job_1744012749162_0002, Tracking URL = http://student-HP-Pro-3330-MT:8088/proxy/
210 Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1744012749162_0002
211 Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
212 2025-04-07 13:52:46,104 Stage-1 map = 0%, reduce = 0%
213 2025-04-07 13:52:57,080 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.44 sec
214 2025-04-07 13:53:11,504 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.72 sec
215 MapReduce Total cumulative CPU time: 2 seconds 720 msec
216 Ended Job = job_1744012749162_0002
217 Loading data to table new_jyoti.new_jyoti.emp_new_idx.emp_id_
218 MapReduce Jobs Launched:
219 Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.72 sec HDFS Read: 10008 HDFS Write: 477 SUCCESS
220 Total MapReduce CPU Time Spent: 2 seconds 720 msec
221 OK
222 Time taken: 44.081 seconds
223 hive> show tables;
224 OK
225 emp_new
226 new_jyoti.emp_new_idx.emp_id_
227 new_jyoti.emp_new_idx.emp_name_
228 Time taken: 0.103 seconds, Fetched: 3 row(s)
229 hive> select * from new_jyoti.emp_new_idx.emp_id_;
230 OK
231 NULL    hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [31]
```

```
230 OK
231 NULL    hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [31]
232 1        hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [0]
233 2        hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [8]
234 3        hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [14]
235 4        hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [21]
236 Time taken: 0.221 seconds, Fetched: 5 row(s)
237 hive> create index idx_emp_id on table emp_new(id) as
'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler' with deferred rebuild;
238 OK
239 Time taken: 0.317 seconds
240 hive> alter index idx_emp_id on emp_new rebuild;
241 WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider
using a different execution engine (i.e. tez, spark) or using Hive 1.X releases.
242 Query ID = hduser_20250407135230_0084ceda-3546-4eef-aed2-d408f4d08916
243 Total jobs = 1
244 Launching Job 1 out of 1
245 Number of reduce tasks not specified. Estimated from input data size: 1
246 In order to change the average load for a reducer (in bytes):
247   set hive.exec.reducers.bytes.per.reducer=<number>
248 In order to limit the maximum number of reducers:
249   set hive.exec.reducers.max=<number>
250 In order to set a constant number of reducers:
251   set mapreduce.job.reduces=<number>
252 Starting Job = job_1744012749162_0002, Tracking URL = http://student-HP-Pro-3330-MT:8088/proxy/
application_1744012749162_0002/
253 Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1744012749162_0002
```

```
Desktop) - gedit 2:12 PM
File Open Save Undo Redo Cut Copy Paste Find Replace Search
f1.txt x employee.txt x emp.txt x prac3 x
application_1744012749162_0002/
253 Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1744012749162_0002
254 Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
255 2025-04-07 13:52:46,104 Stage-1 map = 0%, reduce = 0%
256 2025-04-07 13:52:57,080 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.44 sec
257 2025-04-07 13:53:11,504 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.72 sec
258 MapReduce Total cumulative CPU time: 2 seconds 720 msec
259 Ended Job = job_1744012749162_0002
260 Loading data to table new_jyoti.new_jyoti.emp_new_idx.emp_id_
261 MapReduce Jobs Launched: 1
262 Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.72 sec HDFS Read: 10008 HDFS Write: 477 SUCCESS
263 Total MapReduce CPU Time Spent: 2 seconds 720 msec
264 OK
265 Time taken: 44.081 seconds
266 hive> show tables;
267 OK
268 emp_new
269 new_jyoti.emp_new_idx.emp_id_
270 new_jyoti.emp_new_idx.emp_name_
271 Time taken: 0.103 seconds, Fetched: 3 row(s)
272 hive> select * from new_jyoti.emp_new_idx.emp_id_;
273 OK
274 NULL    hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [31]
275 1      hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [0]
276 2      hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [8]
277 3      hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [14]
278 4      hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [21]
```

Plain Text Tab Width: 8 Ln 298 Col 1 INS

HP v185e

```
top) - gedit
Open Save Undo Redo Cut Copy Paste Find Replace Find in Files Go To View Insert Tools Help
.txt employee.txt emp.txt prac3
1 hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [0]
2 hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [8]
3 hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [14]
4 hdfs://localhost:9000/user/hive/warehouse/new_jyoti.db/emp_new/emp.txt [21]
5 Time taken: 0.221 seconds, Fetched: 5 row(s)
6 hive>
7 exit;
8 hduser@student-HP-Pro-3330-MT:/usr/local/hive/bin$ stop-all.sh
9 This script is Deprecated. Instead use stop-dfs.sh and stop-yarn.sh
10 25/04/07 13:58:40 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...
11 using builtin-java classes where applicable
12 Stopping namenodes on [localhost]
13 localhost: stopping namenode
14 localhost: stopping datanode
15 Stopping secondary namenodes [0.0.0.0]
16 0.0.0.0: stopping secondarynamenode
17 25/04/07 13:59:01 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...
18 using builtin-java classes where applicable
19 stopping yarn daemons
20 stopping resourcemanager
21 localhost: stopping nodemanager
22 localhost: nodemanager did not stop gracefully after 5 seconds: killing with kill -9
23 no proxyserver to stop
24 hduser@student-HP-Pro-3330-MT:/usr/local/hive/bin$ ^C
25 hduser@student-HP-Pro-3330-MT:/usr/local/hive/bin$
```

Plain Text Top Width: 8 Line: 298 Col: 1 INS



Auto Menu

HP v185e